

OpenBLAS physics kernels: throughput vs problem size
double precision, 1 core — L2 peak at $n \approx 10^2 - 10^3$, DRAM plateau from $n \approx 10^5$

